

[FOUNDER TOOLKIT] GO-TO-MARKET TOOLKIT

From one sale / to a system.

Fit gets you your first customers. A **repeatable go-to-market motion** turns them into a business — and it's the exact thing a Series A underwrites. This is how to define who you sell to, how you reach them, and how selling stops depending on the founder.

HOW TO READ THIS TOOLKIT

Three layers / one revenue engine.

Go-to-market is not marketing campaigns — it’s the whole system for how your product reaches and keeps customers. Get the **target and message** right first, then choose **how you reach and charge** them, then turn founder-led selling into a **repeatable motion**. Most GTM fails because founders skip the target and jump straight to channels.

LAYER 1	Who & why	Aim the engine Target & message	Ideal customer profile · positioning · the GTM motion
LAYER 2	How you reach & charge	Build the channels Distribution & price	Channel strategy · pricing & packaging · demand & pipeline
LAYER 3	Make it repeatable	Scale the motion Founder-led to engine	Founder-led sales · the first sales hire · the repeatable engine

LAYER 1 • WHO & WHY

Aim / the engine.

Before a single channel or campaign, decide exactly who you're selling to and why they'd choose you over the alternative. The most common GTM failure is jumping to "how do we reach people" before answering "who, and why us."

Ideal customer profile

Positioning

The GTM motion



TOOL 1.1 · IDEAL CUSTOMER PROFILE

Ideal / customer profile

Your ICP is the specific kind of customer you can win, serve, and keep best — sharpened by the fit work you've already done. Get hyper-specific; trying to sell to everyone sells to no one.

— WHAT IT IS

An **ideal customer profile** describes the customer you're built to win: their role, company type, trigger to buy, and the pain you uniquely solve. It's the natural extension of your validated beachhead — narrowed further by who actually *converts and retains*, not just who has the problem. The discipline is ruthless focus: a tight ICP makes every later choice (channel, message, price) sharper, because you're optimising for one buyer rather than averaging across many. Define the **economic buyer**, the **champion**, and any **gatekeeper** — in B2B these are rarely the same person.

— THE QUESTIONS IT ANSWERS

- ◇ Exactly who converts and retains best — not just 'has the pain'?
- ◇ Who's the **economic buyer** vs. the **champion** vs. the **gatekeeper**?
- ◇ What **trigger** makes them buy now?
- ◇ Are we narrow enough to win, or trying to serve everyone?

— WHAT "GOOD" LOOKS LIKE

- ✓ A specific, written ICP — role, company, trigger
- ✓ Buyer, champion, and gatekeeper named separately
- ✓ Built from who actually converts, not just who has the pain
- ✓ Narrow enough that messaging can be specific

— WHAT NOT TO DO — AND WHY

Defining the ICP as 'anyone with the problem.' A broad target feels like a bigger market but produces vague messaging, scattered channels, and a sales process that fits no one well. "Enterprises" or "insurers" isn't an ICP. The other trap is confusing the *user* (who feels

— ONE ACCOUNT, THREE ROLES

CHAMPION	Feels the pain, sells internally
BUYER	Holds the budget, signs
GATEKEEPER	Can veto (security, legal, DPO)

In B2B these are rarely the same person.

■ AEGIS E0

Aegis's fit work had already produced a precise ICP, and GTM sharpened it further: **user/champion = Head of Underwriting Data** at a top-10 DACH P&C carrier; **economic buyer = Group IT** (the <€50k authority); **gatekeeper = the DPO** on PII. Three distinct people, one account — naming them separately is what made the sales motion work, because each needed a different conversation.

TOOL 1.2 · POSITIONING

Positioning

Positioning is the context that makes your value obvious. The key question isn't “what are our features” — it's “what would the customer use if we didn't exist,” because that alternative defines how they judge you.

— WHAT IT IS

April Dunford's method: positioning is the **frame of reference** you put your product in so its value is obvious to the right customer. Build it from the customer's view, not a slogan. Start with the **true competitive alternative** — what they'd actually do if you didn't exist (often a spreadsheet, a manual process, or nothing), not a list of rival startups. Then identify the **unique attributes** you have that the alternative lacks, the **value** those enable, the **customers** who care most about that value, and the **market category** that makes all of it click. Position against the real alternative, not an imagined competitor set.

— THE QUESTIONS IT ANSWERS

- ◇ *If we didn't exist, what would the customer **actually use**?*
- ◇ *What can we do that the **real alternative** can't?*
- ◇ *What **value** does that unlock — in the customer's terms?*
- ◇ *What **market category** makes our value immediately obvious?*

— WHAT "GOOD" LOOKS LIKE

- ✓ Positioning built from the true competitive alternative
- ✓ Unique attributes tied to real, valued outcomes
- ✓ A category that frames the product for the buyer
- ✓ Message a customer repeats correctly back to you

— WHAT NOT TO DO — AND WHY

Positioning against startups no one's heard of — or with a made-up slogan.

Founders list 15 obscure competitors and claim “ease of use,” while customers are actually comparing them to a spreadsheet or doing nothing. If you position against the wrong alternative, your whole pitch misses. And a positioning *statement* with blanks you simply

Patrick Connolly

— POSITIONING · DUNFORD

- 1 True competitive alternative
- 2 Unique attributes it lacks
- 3 Value those unlock
- 4 Customers who care most
- 5 Market category that frames it

■ AEGIS EO

Aegis positioned against the **true alternative** — not other startups (there were none doing integration), but carriers' own **stalled in-house attempts and unused Maxar subscriptions**. The unique attribute the alternative lacked: **PII-compliant integration that actually clears the DPO**. The category that made it obvious: not “another data vendor” but “the workflow layer for EO data you already pay for.”

Founder Toolkit

TOOL 1.3 · THE GTM MOTION

The GTM / motion

How you sell should match how your customer buys. The deal size and complexity largely decide whether you lead with the product, with salespeople, or with targeted account campaigns.

— WHAT IT IS

Pick the motion that fits your economics and buyer. **Product-led (PLG)** works when the deal is small (ACV under ~\$5k), time-to-value is fast, and an individual can decide — the product sells itself via free trial or freemium. **Sales-led (SLG)** fits complex, higher-value deals (ACV above ~\$25–50k) with multiple stakeholders that need consultative selling. **Account-based (ABM)** suits a defined list of 50–500 high-value targets (ACV \$100k+) with long cycles. Most companies start with one and layer in another as they scale; the motion must mirror the customer's actual buying process.

— THE QUESTIONS IT ANSWERS

- ◇ What's our **ACV** and deal complexity — and which motion fits?
- ◇ Can an individual decide (PLG), or are there **many stakeholders** (SLG/ABM)?
- ◇ How does our customer **actually buy** today?
- ◇ Are we forcing a motion that fights the buying process?

— WHAT "GOOD" LOOKS LIKE

- ✓ Motion matched to ACV and buying complexity
- ✓ One primary motion, not three half-built ones
- ✓ The sales process mirrors how customers buy
- ✓ A clear reason for the motion, not a default

— WHAT NOT TO DO — AND WHY

Choosing the motion by fashion, not fit. Bolting self-serve PLG onto a \$100k product with a six-month committee sale (or hiring an expensive sales team for a \$20/month tool) wastes money and fights the buyer. Running all three motions at once with limited

— MOTION BY DEAL SIZE

<\$5K ACV	PLG — product sells itself
\$25K+ ACV	SLG — consultative selling
\$100K+ ACV	ABM — 50–500 named accounts

Start with one. Layer later.

■ AEGIS E0

Aegis's economics dictated the motion: ~€30–50k ACV, a multi-stakeholder carrier sale (champion + Group IT + DPO), long evaluation. That's squarely **sales-led, trending account-based** — a defined list of ~10 DACH carriers, consultative selling, founder-led at first. PLG would have been absurd for a product bought by a procurement committee. The motion mirrored exactly how carriers actually buy.

LAYER 2 · HOW YOU REACH & CHARGE

Build / the channels.

With the target and message set, decide how you'll reach customers and what you'll charge. Distribution kills more good products than competitors do — and pricing sends a signal as loud as it collects revenue.

Channel strategy

Pricing & packaging

Demand & pipeline

2

TOOL 2.1 · CHANNEL STRATEGY

Channel / strategy

A great product with no distribution is a hobby. Most startups have a product bias and neglect channels — so test several deliberately, then concentrate on the one or two that actually work.

— WHAT IT IS

A channel is how you reliably reach and acquire customers — content, outbound, paid, partnerships, communities, events, and more. Founders systematically **underrate distribution** and assume a good product will sell itself; it won't. Treat channels as **experiments**: test a handful cheaply against cost-to-acquire vs. lifetime value, then **concentrate** on the one or two that produce predictable results rather than spreading thin. The goal by Series A is to have **narrowed to one or two channels that work repeatably** — not a dozen you dabble in.

— THE QUESTIONS IT ANSWERS

- ◇ Which channels could plausibly reach our ICP — have we listed them?
- ◇ Are we **testing** several cheaply, or betting on one untested?
- ◇ What's the **CAC vs. LTV** on each channel we've tried?
- ◇ Have we **concentrated** on the one or two that work?

— WHAT "GOOD" LOOKS LIKE

- ✓ Channels treated as experiments, measured on CAC:LTV
- ✓ A deliberate test across several, not a single guess
- ✓ Concentrated on the 1–2 that produce predictable results
- ✓ Distribution taken as seriously as product

— WHAT NOT TO DO — AND WHY

Assuming the product will sell itself — or dabbling in ten channels. Product bias is the classic founder blind spot: pour everything into building, treat distribution as an afterthought, and watch a great product die unbought. The opposite failure is spreading thin across every channel so none gets a real test. Test deliberately, kill what doesn't work, and double down on what does — distribution is a discipline, not luck.

Patrick Connolly

— TEST, THEN CONCENTRATE

Test several channels → Measure CAC : LTV →

Concentrate on the 1–2 that work

Distribution kills more good products than competitors do.

■ AEGIS E0

Aegis didn't spray channels — its ICP lived in a tiny, reachable world. The channels that worked were **warm introductions and referrals** (Allianz referred two peer carriers unprompted; AXA opened its Group IT) plus targeted industry presence. For ~10 named DACH carriers, paid ads would have been pure waste. **One or two channels, matched to where the buyers actually are.**

TOOL 2.2 · PRICING & PACKAGING

Pricing / & packaging

Price on the value you deliver, not your costs. Pricing is a signal as much as a number — too low can broadcast low quality, especially in high-stakes B2B — and packaging shapes who buys what.

— WHAT IT IS

Set price by **value** (what the outcome is worth to the customer), not cost-plus (your expenses plus a margin), which systematically under-prices real value. Anchor against the customer's alternative: if a consultant costs €50k, software that does the job for €12k is an easy yes *and* can command more than its build cost. Price is also a **signal** — in high-compliance, high-stakes B2B, pricing too low implies low quality and can repel the buyer. **Packaging** (tiers, what's included) steers buyers to the right plan and creates expansion room. Test willingness to pay in discovery, not after launch.

— THE QUESTIONS IT ANSWERS

- ◇ Are we pricing on **value** delivered, or cost-plus?
- ◇ What's the customer's **alternative**, and what does it cost them?
- ◇ Does our price **signal** the right quality for this buyer?
- ◇ Does packaging steer buyers and leave room to **expand**?

— WHAT "GOOD" LOOKS LIKE

- ✓ Value-based pricing anchored to the alternative's cost
- ✓ Price level sends the right quality signal
- ✓ Packaging guides buyers and enables expansion
- ✓ Willingness to pay tested with real buyers

— WHAT NOT TO DO — AND WHY

Cost-plus pricing — or pricing low to win. Cost-plus ignores the value you create and leaves money on the table. Undercutting to attract customers feels safe but can backfire: in high-risk regulatory or enterprise contexts, a suspiciously low price signals “cheap and

— VALUE, NOT COST-PLUS

COST-PLUS Expenses + margin — under-prices value

VALUE-BASED **Anchored to the alternative's cost**

Price is a signal: too low =
“low quality” in high-stakes B2B.

■ AEGIS E0

Aegis priced against the alternative, not its costs. Carriers' status quo — stalled in-house builds and consultant engagements — ran far higher than an integration subscription, so a **value-anchored €30–50k/year** was both an easy ROI case *and* a credible quality signal to Group IT. Pricing it cheap would have implied it couldn't be trusted with **PII-grade compliance** — the opposite of the message.

TOOL 2.3 · DEMAND & PIPELINE

Demand / & pipeline

Revenue starts as a pipeline: a tracked path from “never heard of you” to “signed.” Make the funnel visible, measure conversion at each stage, and you can diagnose and forecast instead of guessing.

— WHAT IT IS

Demand generation creates awareness and interest; the **pipeline** is the tracked path from first contact to closed deal — stages like lead → qualified → demo → proposal → closed. Making it explicit lets you measure **conversion at each stage**, find the bottleneck, and forecast revenue rather than hope for it. Early on this is a simple, honest tracker, not a CRM ritual. Watch **pipeline coverage** (enough live deals to hit the target given your win rate) and **velocity** (how fast deals move). A real pipeline reflects genuine buyer intent — not a wish-list of logos.

— THE QUESTIONS IT ANSWERS

- ◇ What are our **pipeline stages**, first contact to close?
- ◇ What's our **conversion** at each stage — where's the bottleneck?
- ◇ Do we have enough **coverage** to hit the target at our win rate?
- ◇ Does the pipeline reflect real **buyer intent**, or hope?

— WHAT "GOOD" LOOKS LIKE


- ✓ A visible funnel with stage-by-stage conversion
- ✓ The bottleneck stage identified and worked
- ✓ Pipeline coverage sufficient for the target
- ✓ Deals reflect real intent, not logo wish-lists

— WHAT NOT TO DO — AND WHY

A ‘pipeline’ that’s a list of dream logos. Tracking accounts you’d love rather than deals with real, evidenced intent produces a forecast that’s pure fiction — and hides the truth that nothing is actually moving. The other trap is no pipeline at all: selling by vibes, unable to say why deals stall or how many you’ll close. A pipeline is a measurement instrument; if it doesn’t reflect reality, it misleads worse than nothing.

Patrick Connolly

— THE PIPELINE



Measure conversion at each stage.
Watch coverage and velocity.

■ AEGIS E0

Even selling founder-led to a handful of carriers, Aegis tracked an honest pipeline: which of the ~10 DACH targets were at intro, pilot discussion, MOU, or paid. Munich Re moved intro → **signed pilot MOU**; AXA sat at **Group-IT-intro**; Allianz’s two referrals entered as fresh leads. Stage-by-stage, it could see exactly where each deal was — and what to do next — instead of guessing.

LAYER 3 • MAKE IT REPEATABLE

Scale / the motion.

Founder-led sales is the right way to start — and a trap if it never ends. The Series A bar is a motion that works when the founder isn't in the room. This layer is the transition from heroic selling to a repeatable engine.

Founder-led sales

The first sales hire

The repeatable engine

3

TOOL 3.1 · FOUNDER-LED SALES

Founder-led / sales

In the early days the founder must sell — no one understands the product or customer better. The goal isn't to sell forever; it's to learn the sales process so well that it can be written down and handed off.

— WHAT IT IS

Early selling should be **founder-led**: the founder runs the conversations, adapts the pitch in real time, and hears every objection first-hand. This isn't just closing deals — it's **discovering the sales process**: the language buyers use, the objections that recur, the triggers that turn interest into a signature. Each conversation refines product, message, and motion. The objective is to extract the **repeatable pattern** from these wins — what consistently works — so it can later be documented and taught. Founders should sell until the pattern is clear, not delegate before it exists.

— THE QUESTIONS IT ANSWERS

- ◇ Is the **founder** running the early sales conversations?
- ◇ Are we capturing the **objections and language** that recur?
- ◇ What's the **pattern** in the deals we win?
- ◇ Could we **write down** what consistently works yet?

— WHAT "GOOD" LOOKS LIKE

- ✓ Founder selling directly and learning the process
- ✓ Recurring objections and winning language captured
- ✓ A discernible pattern across closed deals
- ✓ Insight feeding back into product and message

— WHAT NOT TO DO — AND WHY

Handing off sales before the process exists. Hiring a salesperson to “do sales” while the motion is still undiscovered fails reliably — they have nothing repeatable to execute, can't adapt the pitch the way a founder can, and miss the product feedback the founder would have caught. The mirror error is the founder clinging to every deal forever, becoming the bottleneck. Sell personally until the pattern is clear — then, and only then, hand it off.

Patrick Connolly

— FOUNDER-LED SALES = R&D

Founder sells → Hears objections, adapts →

Extracts the repeatable pattern

Sell until the process is clear — then write it down.

■ AEGIS E0

Aegis's founders ran every early carrier conversation themselves — and that's how they learned the **real sales pattern**: the champion (underwriting data) had to be armed to sell Group IT internally, and **nothing closed until the DPO's PII concern was answered first**. No hired rep would have surfaced that. They sold founder-led precisely to extract the motion before scaling it.

TOOL 3.2 · THE FIRST SALES HIRE

The first / sales hire

The first salesperson should validate that someone other than the founder can close — not be sent to discover a motion that doesn't exist yet. Hire once you have a documented, repeatable process.

— WHAT IT IS

Make the first sales hire when the founder has **closed ~10 customers the same way**, the process is **documented**, the ICP and message are clear, and there's consistent pipeline. The hire's job is to prove the motion **transfers** — that a non-founder can run the playbook and close. This is the precise proof point Series A investors want: **non-founder-led deals as a meaningful share of revenue**. Hire the rep, give them the documented playbook and warm pipeline, and watch whether they ramp; once one rep succeeds, you have a motion worth adding a second to.

— THE QUESTIONS IT ANSWERS

- ◇ Has the founder closed **~10 deals the same way**?
- ◇ Is the sales process **documented** enough to hand over?
- ◇ Are we hiring to **prove transfer** — or to outsource discovery?
- ◇ Can we show **non-founder-led** deals closing?

— WHAT "GOOD" LOOKS LIKE

- ✓ First rep hired after a repeatable process exists
- ✓ A documented playbook + warm pipeline to inherit
- ✓ The test is whether the motion transfers off the founder
- ✓ Non-founder-led revenue — the Series A proof point

— WHAT NOT TO DO — AND WHY

Hiring a 'VP of Sales' to figure it out for you. Bringing in an expensive senior seller before the motion is proven is a top reason early sales hires fail — they're handed a blank page, not a playbook, and the founder's tacit knowledge never transfers. Equally, never proving a non-founder can close leaves you founder-dependent, which caps the company and fails diligence. Document first, hire to transfer, then scale.

Patrick Connolly

— WHEN TO HIRE THE FIRST REP

- | | |
|---------|------------------------------------|
| READY | ~10 closed + documented + pipeline |
| NOT YET | No repeatable process to hand over |

The hire's job: prove the motion transfers off the founder.

■ AEGIS E0

Aegis sequenced it deliberately. The plan: only after the founders had converted the Munich Re pilot and closed two more DACH carriers **the same way** — a documented motion — would they hire a first commercial rep, handing over the playbook and the warm Allianz referrals. The goal was explicit: prove a non-founder can close, because **that's exactly what the Series A would underwrite**.

TOOL 3.3 · THE REPEATABLE ENGINE

The repeatable / engine

A go-to-market engine is repeatable when you can predict what input produces what output: spend here, get pipeline there, convert at this rate. That predictability — not raw revenue — is what earns the next round.

— WHAT IT IS

The destination is a **repeatable, predictable** motion: a defined ICP, a message that lands, one or two working channels, a documented sales process a non-founder can run, and known conversion rates that let you forecast. At that point growth becomes an input-output system — add a rep or a channel dollar and you can predict the revenue. This is precisely the Series A bar: **real, repeatable, scalable** revenue with healthy unit economics (the metrics from the Business Operations and Fundraising toolkits — LTV:CAC, CAC payback, retention). Don't scale spend until the engine is proven; then pour fuel on what works.

— THE QUESTIONS IT ANSWERS

- ◇ Can we **predict** revenue from a given input (rep, channel spend)?
- ◇ Is the motion documented and run by **non-founders**?
- ◇ Are the **unit economics** healthy (LTV:CAC, payback, retention)?
- ◇ Are we **scaling spend** only on a proven engine?

— WHAT "GOOD" LOOKS LIKE

- ✓ Predictable input → output across the funnel
- ✓ A documented motion that runs without the founder
- ✓ Healthy, improving unit economics
- ✓ Spend scaled only after the engine is proven

— WHAT NOT TO DO — AND WHY

Pouring fuel on an unproven engine. Raising and then spending hard on sales and marketing before the motion is repeatable just burns cash faster — the premature-scaling trap that kills more funded startups than almost anything. Predictability comes first; scale

Patrick Connolly

— REPEATABLE = PREDICTABLE

Input (rep / spend) → Known conversion → Predictable revenue

Series A funds scaling a proven engine — not finding one.

■ AEGIS E0

Aegis's Series A story was the engine, not just the revenue: a defined DACH-carrier ICP, PII-compliant positioning, a working referral+pilot channel, a documented sell-through-the-champion-to-Group-IT motion, and a first rep proving it transferred. With **predictable conversion and healthy economics**, scaling spend became a calculated bet — the exact thing the round was raised to do.

Founder Toolkit

PUTTING IT TOGETHER

Build the engine / before you scale it.

Go-to-market is the bridge between a product people want and a business that grows. Aim it (who, and why you) before you build channels; choose how you reach and charge to match how customers buy; then turn founder-led selling into a motion that runs without you. A repeatable engine — not raw revenue — is what earns the Series A.

Layer 1

Aim

A precise ICP, positioning against the real alternative, and a motion matched to how customers buy.

Layer 2

Reach & charge

Test channels and concentrate on what works; price on value; make the pipeline visible and honest.

Layer 3

Repeat

Sell founder-led to learn the motion, hire to prove it transfers, then scale a predictable engine.

The point

Predictable

When input reliably produces output, you've built the engine Series A pays to scale.

THE EVIDENCE BASE

Sources / & further reading.

This toolkit favours respected operator sources (Tunguz, Dunford, Thiel, Moore) and current Series-A market data. The “what not to do” guidance reflects documented failure patterns, not opinion. It connects directly to the PMF toolkit (the beachhead becomes the ICP) and to Business Operations & Fundraising (the unit economics). Aegis EO is a fictional teaching venture; its motion is illustrative.

The five GTM elements

Tomasz Tunguz (Theory Ventures) and GTM consensus: target customer, value proposition, channel, pricing/packaging, sales motion — aligned around a repeatable acquisition engine.

PLG vs. SLG vs. ABM

Tomasz Tunguz; industry practice: motion by ACV — PLG <~\$5k, sales-led >~\$25–50k, account-based for \$100k+ named accounts.

Channels as experiments

Gabriel Weinberg & Justin Mares (*Traction*): test channels broadly, then concentrate on the one or two that work (the Bullseye idea).

Founder-led sales → first hire

Tomasz Tunguz; YC: founder sells until ~10 closed the same way and the process is documented, then hires to prove the motion transfers.

Beachhead → expand

Geoffrey Moore (*Crossing the Chasm*): win a narrow segment, earn references, expand — the ICP’s lineage from the PMF beachhead.

Obviously Awesome (2019)

April Dunford. Positioning from the true competitive alternative → unique attributes → value → best-fit customers → market category.

Distribution

Peter Thiel (*Zero to One*), Marc Andreessen: founders systematically underrate distribution; a great product with no channel fails.

Value-based pricing

Pricing practice: price on value and the customer’s alternative, not cost-plus; price as a quality signal in high-stakes B2B.

Series A = repeatable GTM

CRV, a16z, Carta (2026): Series A funds scaling a codified, repeatable go-to-market motion — non-founder-led deals as the key proof point.

Unit economics

David Skok; consensus: LTV:CAC ≥ 3 , CAC payback <12mo, retention — the economics that make an engine worth scaling (see Business Ops & Fundraising toolkits).