PATRICK
CONNOLLY

[FOUNDER TOOLKIT] THE PILOT PLAYBOOK

The descent / and the landing.

You proved the pain is real. Now you must prove you can fix it for **one** customer — without burning your runway, and in a way that hands you research, a reference, and the next deal. This is how to land your first real pilot and take off with more than you arrived with.

June 2026

Developed for founders I mentor at the International Space University and ISU Incubator

HOW TO READ THIS PLAYBOOK

Three phases / one approach plate.

A pilot is one maneuver, not the whole flight. It lives in the transition from **validated problem** to **solution signal** — PMF told you the bet is right; this lands your first real customer engagement. Like an approach plate, each phase assumes the one above it is cleared: set up the **approach**, run the **anding** without crashing, then harvest the **next runway**. Don't skip ahead.

PHASE 1	**Cleared to descend?**	**Approach** Set up the landing	Readiness gate · **design-partner selection** · the decision-making unit
PHASE 2	**Without crashing**	**Landing** Run the pilot	Scope the PoC · **success & kill criteria** · pitch it & derisk cheap
PHASE 3	**Take off again**	**Next runway** Harvest three things	Capture & read the verdict · **the three harvests** · set up customer #2

LAYER 1 · CLEARED TO DESCEND?

Approach / set up the landing.

Most pilots fail here, silently, before they start: wrong customer, fuzzy scope, no one who can actually say yes. The approach decides everything downstream — confirm you're ready, pick the one customer, and map who really decides.

Readiness gate

Design-partner selection

The decision-making unit

1

CHECKPOINT 01 · READINESS GATE

The readiness / gate

Are you actually ready to pilot — or still validating? A pilot tests a solution. If you're still discovering the problem, a pilot gives you noisy, expensive, misleading data. Clear this gate first.

— WHAT IT IS

A pilot is a test of a **solution**, which only produces clean signal once the **problem** is genuinely validated (the PMF toolkit's job). Before you descend, confirm four things: you can state the problem in one sentence with a named role, and strangers have confirmed that exact pain **unprompted**; you have a solution hypothesis concrete enough to deliver in **weeks**, not a vision; there's real **pull** (a customer asked "when can I try it?" rather than you pushing); and you can afford to run the pilot **and survive it failing**. Pull, not your own confidence, is what clears this gate.

— THE QUESTIONS IT ANSWERS

- *Can I state the problem in one sentence, with a **named role**?*
- *Has that exact pain been confirmed by strangers, **unprompted**?*
- *Is there real **pull** — did a customer ask to try it?*
- *Can we afford to run this pilot **and survive it failing**?*

— WHAT "GOOD" LOOKS LIKE

- ✓ A falsifiable solution hypothesis with a metric attached
- ✓ Inbound pull from ≥1 named customer
- ✓ You can describe a "win" in numbers
- ✓ The problem survived ~40+ conversations

— WHAT NOT TO DO — AND WHY

Piloting to avoid talking to customers. A pilot feels like progress, so teams jump to it to escape the discomfort of discovery. If you can't yet name the problem and the buyer, you're not derisking — you're building on sand with a logo attached. The expensive, misleading data from a premature pilot is worse than no data, because it feels authoritative. Validate the problem first; then descend.

— CLEARED TO DESCEND?

PROBLEM	One sentence, named role, confirmed unprompted
SOLUTION	Concrete, deliverable in weeks
PULL	A customer asked to try it
SURVIVE	Can afford it failing

AEGIS EO

After 52 conversations, Aegis had killed three assumptions and landed a final hypothesis: carriers already pay for imagery; the bottleneck is **PII-compliant workflow integration**. The pull was unmistakable — Munich Re signed a free pilot MOU before the sprint even ended. That **inbound, not the team's confidence**, is what cleared this gate.

CHECKPOINT 02 · DESIGN-PARTNER SELECTION

Choosing / the runway

Pick the one customer by fit, not by who's loudest. Your first pilot customer is a design partner, not just revenue — the wrong one teaches the wrong lessons and produces a case study no one else recognizes.

— WHAT IT IS

Score candidates on five criteria. **Acute pain**: they feel this badly enough to make time for you. **Representative**: they look like your next 20 customers, not an exotic edge case. **Reachable decision**: the people who can say yes will actually be in the room. **Reference-able**: their logo and a quote would move customer #2. **Survivable**: they won't demand so much custom work that you accidentally build a consultancy. Choose on **written criteria**, not enthusiasm — the loudest customer is rarely the most representative.

— THE QUESTIONS IT ANSWERS

- *Do they feel the pain **acutely** enough to make time?*
- *Are they **representative** of the next 20, not an edge case?*
- *Will the people who can **say yes** be in the room?*
- *Would their logo and quote **move customer #2**?*

— WHAT "GOOD" LOOKS LIKE

- ✓ One named account, chosen on written criteria
- ✓ A "why them first" you can defend
- ✓ They're in your beachhead, not adjacent
- ✓ A credible #2 and #3 already visible behind them

— WHAT NOT TO DO — AND WHY

The enthusiastic outlier. The most eager customer is often the least representative — a power user with a niche need and a budget no one else has. You'll win them, build for them, and discover the win doesn't transfer to anyone else. Picking by loudness instead of fit produces a bespoke success that teaches the wrong lessons and a case study the rest of your market doesn't recognize.

— SCORE ON FIVE CRITERIA

PAIN	Acute enough to make time
REPRESENTATIVE	Looks like the next 20
REACHABLE	Deciders in the room
REFERENCE-ABLE	Logo + quote moves #2
SURVIVABLE	Won't demand a consultancy

AEGIS EO

Aegis chose Munich Re over an eager UK broker. Why: Munich Re sits dead-center in the **DACH beachhead**, a warm alumni intro put the right people in the room, and Allianz had already referred two peer carriers — proof the win would transfer. The broker was louder; **Munich Re was representative**. Fit beat enthusiasm.

CHECKPOINT 03 · THE DECISION-MAKING UNIT

Who's / in the tower

Map who really decides before you scope anything. A pilot dies when an invisible stakeholder vetoes it — find the champion, the economic buyer, the gatekeeper, and the blocker now, not after a no.

— WHAT IT IS

Name four roles in the account, with real people attached. The **champion** feels the pain, wants you to win, and will sell internally for you. The **economic buyer** controls the budget and the real yes. The **gatekeeper** can't say yes but can say no — security, legal, procurement, IT, the DPO. The **blocker** loses something if you succeed: whose status quo you threaten. Map all four before scoping, because the requirement you don't know about on day one becomes the veto that stalls you for nine months at review.

— THE QUESTIONS IT ANSWERS

- *Who is the **champion** — and have I armed them to sell up?*
- *Who is the **economic buyer** who controls the real yes?*
- *Who is the **gatekeeper** — and what do they require, on day one?*
- *Who is the **blocker**, and what do they fear?*

— WHAT "GOOD" LOOKS LIKE

- ✓ Each of the four roles has a real name
- ✓ The champion has ammunition to sell upward
- ✓ The gatekeeper's requirements known on day one
- ✓ You know what the blocker fears, and have an answer

— WHAT NOT TO DO — AND WHY

Selling to your champion and calling it done. The champion's enthusiasm is necessary but not sufficient — the deal that "everyone loved" stalls for nine months in a procurement or compliance review you never mapped. An invisible gatekeeper (security, legal, a DPO) can veto a pilot the champion already said yes to. Map the whole decision-making unit before you scope, so the no comes early and cheap, not late and fatal.

— THE FOUR ROLES

CHAMPION	Feels the pain, sells internally
BUYER	Controls budget & the real yes
GATEKEEPER	Can't say yes, can say no
BLOCKER	Loses if you win

AEGIS EO

Aegis named all four. Champion: the **Head of Underwriting Data**, who felt the pain daily. Economic buyer: **Group IT** — 11 of 13 buyers confirmed authority sat there, not at the desk. Gatekeeper: the **DPO**, because PII compliance was the unspoken wall that killed every in-house attempt. Naming the DPO early is why Aegis scoped compliance in from the start instead of discovering it at review.

LAYER 2 · WITHOUT CRASHING

Landing / run the pilot.

Now you execute: scope the narrowest test that proves the most, pre-commit what counts as a win, sell it as a low-risk experiment, and prove value with effort rather than a finished product. The discipline here is to learn the most for the least.

Scope the PoC

Success & kill criteria

Pitch it & derisk cheap

2

CHECKPOINT 04 · SCOPING THE POC

The smallest test / that proves the most

Scope is where pilots quietly become death marches. Every "while we're at it" feature adds weeks and dilutes the signal. Design the narrowest engagement that can still produce a clear yes — and write the line down with the customer.

— WHAT IT IS

Draw the scope line explicitly, in writing, with the customer. **In scope** (proves the core bet): the single workflow that carries the most pain; one real, measurable before/after; the minimum to make a champion say "I'd buy this"; and whatever the gatekeeper requires to allow it at all. **Out of scope** (name it out loud): integrations beyond the one that matters, edge cases and polish, anything serving a customer who isn't #1, and scale/performance/multi-user — not yet. The questions that set scope: what is the *one thing* this pilot must prove, and what's the shortest timebox that still yields real data — weeks, not quarters?

— THE QUESTIONS IT ANSWERS

- *What is the **one thing** this pilot must prove?*
- *What's the shortest **timebox** that still produces real data?*
- *What will we **deliberately not do** — said out loud?*
- *Is the line drawn **in writing, with the customer**?*

— WHAT "GOOD" LOOKS LIKE

- ✓ A single workflow that carries the most pain
- ✓ One measurable before/after, timeboxed to weeks
- ✓ Out-of-scope named explicitly, in writing
- ✓ Just enough to make a champion say "I'd buy this"

— WHAT NOT TO DO — AND WHY

Scope creep dressed as helpfulness. Every "while we're at it" addition feels generous and quietly turns a one-month test into a one-quarter death march that dilutes the signal you came for. The mirror failure is scoping so vaguely that no single result can count as a

— DRAW THE LINE · IN WRITING

IN	One workflow · one before/after · gatekeeper's minimum
OUT	Extra integrations · polish · scale · customer #2's needs

Shortest timebox with real data: weeks, not quarters.

AEGIS EO

Aegis scoped to **one workflow at one carrier for one month**: pipe existing Maxar imagery into Munich Re's claims process and measure a single number — time-to-claim. No new data, no multi-carrier support, no polish. The whole pilot existed to move **one metric the buyer already cared about** — everything else was named out of scope, on paper.

CHECKPOINT 05 · SUCCESS & KILL CRITERIA

Decision height / set in advance

Pilots without pre-committed thresholds always "kind of work," because you grade them after seeing the result. Set your decision height now — win, learn, no-go — so the data decides, not your hope.

— WHAT IT IS

Before touching anything, pre-commit three bands. **WIN**: the metric clears the bar you set in advance → convert to paid, trigger the case study, ask for the referral. **LEARN**: mixed — it moves but not enough, and you understand why → one scoped iteration, re-test against the *same* bar. **NO-GO**: the metric misses, or the buyer won't pay even if it works → stop, bank the learning, re-examine the bet (not the build). Write the single success metric and number down, share it, and confirm the customer agrees these are the right measures — or you'll win on yours and still lose the deal.

— THE QUESTIONS IT ANSWERS

- *What single **metric, at what number**, equals success?*
- *What result would make us **walk away** — said out loud?*
- *Does the **customer agree** these are the right measures?*
- *Is willingness-to-pay part of the win bar, not just the metric?*

— WHAT "GOOD" LOOKS LIKE

- ✓ A written win bar, shared with the customer
- ✓ A pre-committed no-go line you'll honor
- ✓ Three bands, each with an automatic next action
- ✓ The customer agrees the measures are the right ones

— WHAT NOT TO DO — AND WHY

Grading on a curve after the fact. Without a pre-set kill line, no pilot ever fails — you'll always find a story for why it "showed promise," and burn another quarter on a bet that already answered no. Setting thresholds *after* seeing the data isn't analysis, it's rationalization. Commit the win, learn, and no-go bands in writing before the data exists, so the result decides your next move — even when the answer is stop.

— PRE-COMMIT THE BANDS

WIN	Clears the bar → convert + case study + referral
LEARN	Moves, not enough → iterate once, re-test
NO-GO	Misses / won't pay → stop, re-examine the bet

AEGIS EO

Aegis pre-committed, in the launch brief, before touching the data: **40% time-to-claim cut = build**; 20–40% = iterate once and re-test; under 20% = pivot the bet. Because the bar was set before the result existed, whatever Munich Re's number turned out to be, the team's response was **already honest** — no room to move the goalposts.

CHECKPOINTS 06–07 · PITCH IT & DERISK CHEAP

Low-risk for them / cheap for you

Frame the pilot as their low-risk experiment, not your sales ask — then prove value with effort, not a finished product. Build the product to run the pilot and you've taken the very risk you were trying to avoid.

— WHAT IT IS

The pitch (06): the customer isn't buying a product yet — they're agreeing to spend time, so make yes feel safe, small, and worth it. Include their problem in their words, a tightly bounded ask (this workflow, this long), the one metric you'll move, an explicit exit (if it fails they owe you nothing), and what's in it for them. **Derisk** (07): prove the bet at the lowest cost — **concierge** (do it by hand; they see the outcome), **Wizard-of-Oz** (looks automated, a human runs it), timebox hard, and reuse tools instead of building. Manual is a feature here: it buys the same learning for a fraction of the spend.

— THE QUESTIONS IT ANSWERS

- *Can the customer **restate the deal** in one sentence?*
- *Does the time/cost ask feel **trivial next to the pain**?*
- *Could we **start this week**, with spend capped?*
- *Are we faking the back end (concierge) rather than building it?*

— WHAT "GOOD" LOOKS LIKE

- ✓ Risk visibly sits with you, not the customer
- ✓ The champion can sell the pilot upward
- ✓ A concierge/Wizard-of-Oz MVP, not production code
- ✓ If it fails, you've lost weeks, not the company

— WHAT NOT TO DO — AND WHY

Pitching the 10-year vision — and building the product to run one pilot. Grand TAM slides raise the stakes and the scrutiny: the buyer starts evaluating you as a strategic vendor instead of agreeing to a cheap experiment. Keep the altitude low. And teams routinely spend three months and most of their runway automating a workflow before

— SELL SMALL, DELIVER BY HAND

PITCH	Their words · bounded ask · explicit exit
CONCIERGE	Do it manually — they see the outcome
TIMEBOX	Fixed end date caps spend
REUSE	Stitch tools before writing code

AEGIS EO

Aegis didn't pitch "the EO-insurance platform." They pitched: **"Let us cut your time-to-claim on one workflow for one month, free, using imagery you already pay for."** Low cost, low risk, a number Munich Re cared about — MOU signed in days. Delivery was a **concierge MVP**: run the integration by hand for a month; build the real engineering only if time-to-claim cleared 40%.

LAYER 3 · TAKE OFF AGAIN

Next runway / harvest three things.

One completed pilot should yield three assets for three audiences — R&D, a case study, and your next customer. Most teams take only the revenue and leave the two most valuable, free assets on the table. Read the result honestly first, then harvest all three.

Capture & read the verdict

The three harvests

Set up customer #2

3

CHECKPOINTS 08–09 · CAPTURE & READ THE VERDICT

Measure the before / read it honestly

Capture the baseline now — you can't reconstruct it later — then hold the result against the bar you set, and act on it. This is the wheels-down gate: do what you pre-committed, even if the answer is stop.

— WHAT IT IS

Capture (08): the most valuable thing a pilot produces is a credible before-and-after, and the "before" is perishable — measure the baseline (dated, agreed) on day zero, track the same metric the same way during and after, and log verbatim quotes and friction same-day. **Verdict** (09): compare the outcome to your pre-set thresholds and act automatically. Did the metric clear the win bar — yes or no, not "sort of"? Would the buyer **pay now** (a metric win with no willingness-to-pay is still a no-go)? If LEARN, what's the one change worth another round? If NO-GO, which assumption was wrong?

— THE QUESTIONS IT ANSWERS

- *Did we capture a **dated baseline** both sides agreed?*
- *Did the metric **clear the win bar** — yes or no?*
- *Would the buyer **pay for this now**?*
- *Does the next action follow **automatically** from the band?*

— WHAT "GOOD" LOOKS LIKE

- ✓ A dated baseline both sides signed off on
- ✓ One headline number you can defend to a skeptic
- ✓ Willingness-to-pay confirmed, not assumed
- ✓ The verdict maps cleanly to the pre-set thresholds

— WHAT NOT TO DO — AND WHY

Realizing at the end you never recorded the start — then moving the bar to survive.
Without a baseline captured on day zero, you "know" the pilot worked but can't prove it, and a skeptical customer #2 or investor discounts the whole result. Worse is rewriting success after the fact — the metric came in under target, so suddenly the "real" win was

— CAPTURE, THEN DECIDE

Baseline (day zero, dated) → **Same metric, during/after** → **Verdict vs. pre-set bar**

A metric win with no willingness-to-pay is still a no-go.

AEGIS EO

Because the metric — time-to-claim — was fixed before the pilot, Aegis measured Munich Re's current claim-cycle time as the **baseline on day one**; every later result pointed back to that agreed number. And the pre-set rule did the deciding: clear 40% and convert, 20–40% iterate once, under 20% pivot the bet. The threshold was set before the data existed — so the team's response was **honest by construction**.

CHECKPOINT 10 · THE THREE HARVESTS

Three takeoffs / from one landing

A pilot isn't one output, it's three — R&D, a case study, and a referral — and they serve three different audiences. Harvest all of them deliberately, or you'll leave the two most valuable, free assets on the table.

— WHAT IT IS

Pull three distinct assets from a single completed pilot. **R&D** (for your product): what broke, what they asked for, which feature actually drove the value — this feeds the roadmap and the next build. **Case study** (for marketing): the before/after, the quote, the logo — proof that converts customer #2 faster than any deck you can write. **Referral** (for sales): a warm intro to a peer, asked for *while the champion is happiest* — the cheapest pipeline you will ever generate. Ask "who else should see this?" explicitly, while goodwill is at its peak.

— THE QUESTIONS IT ANSWERS

- *What did this teach us about **what to build next**?*
- *Do we have a **customer-approved case study** — metric, quote, logo?*
- *Have we asked, explicitly, **"who else should see this?"***
- *Are we harvesting all three, not just the revenue?*

— WHAT "GOOD" LOOKS LIKE

- ✓ R&D that reshapes the roadmap
- ✓ A cleared case study: metric + quote + logo
- ✓ A warm referral captured while goodwill peaks
- ✓ All three assets taken deliberately

— WHAT NOT TO DO — AND WHY

Taking only the revenue and moving on. Teams close the pilot, bank the contract, and never write the case study or ask for the intro. The two most valuable, free assets — proof and pipeline — evaporate within weeks as memory fades and the champion gets busy. The window for a glowing quote and a warm referral is widest right after a win; harvest all three then, not "when things settle down."

— THREE ASSETS, THREE AUDIENCES

R&D	For product — what to build next
CASE STUDY	For marketing — proof that converts #2
REFERRAL	For sales — the cheapest pipeline

AEGIS EO

Aegis was set up to harvest all three: **R&D** from running the workflow by hand (what to automate first), a **case study** built on the agreed time-to-claim baseline, and **referrals** already in motion — Allianz had introduced two peer carriers unprompted, and Munich Re's contact offered to co-present at an industry conference. One landing, three takeoffs.

CHECKPOINT 11 · SET UP CUSTOMER #2

One proof point / into a motion

Customer #1 was bespoke and warm. The goal now is to make #2 faster and #3 faster still — turning a single landing into a flight path you can fly again. This is the bridge into the Go-to-Market toolkit.

— WHAT IT IS

Convert one win into a repeatable approach. Repeat the pitch with the **case study where your hypothesis used to be** — proof, not a guess. Standardize what was bespoke for #1, and only re-open what a new customer genuinely changes. Pick a #2 who is **more representative** of the wider beachhead than #1 was, and carry forward the scope, the metric, and the kill criteria as a reusable template. Define the one metric you'll now use to qualify every future pilot. This is where the pilot motion hands off to the GTM toolkit's repeatable engine.

— THE QUESTIONS IT ANSWERS

- *Can we repeat the pitch with **proof** where the hypothesis was?*
- *What was **bespoke** for #1 that we must standardize?*
- *Is #2 **more representative** of the beachhead than #1?*
- *What's the one metric we'll use to **qualify** every pilot?*

— WHAT "GOOD" LOOKS LIKE

- ✓ A reusable pitch built on real proof
- ✓ #2 sourced from #1's referral chain
- ✓ A shorter time-to-pilot than the first
- ✓ A standard scope you tweak, not redraw

— WHAT NOT TO DO — AND WHY

Treating #2 as another snowflake. If every pilot is fully custom, you never escape services and never compound — you're running a consultancy, one bespoke engagement at a time. The discipline is to carry forward the scope, the metric, and the kill criteria as a template, and re-open only what the new customer genuinely changes. Otherwise customer #10 is as slow and as costly as customer #1, and there's no motion to scale.

— ONE LANDING → A FLIGHT PATH

Customer #1 (bespoke, warm) → **Standardize scope + metric + criteria** → **#2 faster, from the referral chain**

Each landing shortens the next.
The motion compounds → GTM toolkit.

AEGIS EO

For Aegis, the referral chain became the runway: with a Munich Re result in hand, the next pitch leads with **proof, not a hypothesis** — and Italian carriers had already said they'd engage once two DACH references existed. Customer #1 was the hard one; #2 was designed to be **faster on purpose**. That repeatable motion is exactly what the Go-to-Market toolkit then scales.

THE LOOP, NOT THE FINISH LINE

A pilot done right / clears you for the next one.

The eleven checkpoints aren't a one-time descent. Each completed pilot feeds the next: sharper scope, a stronger case study, a warmer pipeline. Set up the approach so the customer and scope are right; land it cheaply and read the verdict honestly; harvest R&D, proof, and pipeline. That's how one customer becomes a repeatable motion — and where this playbook hands off to Go-to-Market.

Phase 1

Approach

Ready, partner, decision-making unit. The setup that decides everything downstream — clear each gate before descending.

Phase 2

Landing

Scope the smallest test, pre-commit the kill criteria, pitch low-risk, and derisk with a concierge MVP, not production code.

Phase 3

Next runway

Capture the baseline, read the verdict honestly, then harvest R&D, a case study, and a referral into customer #2.

The point

It compounds

Each landing makes the next faster. The motion compounds; the runway shortens; one win becomes a repeatable engine.

THE EVIDENCE BASE

Sources / & further reading.

This playbook sits between the PMF toolkit (validate the problem) and the Go-to-Market toolkit (scale the motion) — it is the single maneuver of landing your first real customer. It draws on lean-startup and enterprise-sales practice, and on the ISU Flight Map's validated-problem→solution-signal transition. The eleven checkpoints are grouped here into nine for consistency with the toolkit series. Aegis EO is a fictional teaching venture; names, quotes, and figures are illustrative.

Design partners & the concierge MVP

Y Combinator ("Do things that don't scale", Paul Graham); Steve Blank: the first customer is a design partner; fake the back end before you build it.

Pre-committed success criteria

Lean Startup (Ries) / experiment design: set the win, learn, and no-go thresholds before the test, so the data decides, not hindsight.

Measuring the before/after

Standard pilot practice: the baseline is perishable — capture it dated and agreed on day zero, or the case study becomes unprovable.

From one customer to a motion

Geoffrey Moore (bowling-pin); the pilot hands off to the Go-to-Market toolkit's founder-led-to-repeatable engine.

The decision-making unit

Enterprise-sales practice: champion, economic buyer, gatekeeper, blocker — map all four before scoping (connects to the GTM toolkit's ICP).

Concierge & Wizard-of-Oz MVPs

Lean Startup; classic startup practice: prove the bet manually at the lowest cost rather than building production software to run a pilot.

Pivot or persevere

Eric Ries: hold the result against pre-set criteria; a metric win without willingness-to-pay is still a no-go; re-examine the bet, not just the build.

The ISU Flight Map & PMF Sprint

ISU MSS26: the pilot lives in the transition from validated problem to solution signal; this playbook is the maneuver between PMF and GTM.