

[ FOUNDER TOOLKIT ] PRODUCT-MARKET FIT TOOLKIT

# Your venture / is a bet.

Your job before building anything is to find out whether it's the **right** bet — by talking to the people who would actually buy. This is how to frame the bet, run honest discovery, and measure fit with evidence instead of hope.

## HOW TO READ THIS TOOLKIT

# Three layers / one loop.

Product-market fit isn't found by building — it's found by **learning**. Frame a falsifiable bet, run discovery that gets the truth instead of politeness, then measure fit honestly and decide. The three layers form one engine: aim, test, learn, re-aim. Skip the framing and your conversations have nothing to test; skip the discipline and you hear only what you hoped to.

LAYER 1

## Aim

### Frame the bet

A falsifiable hypothesis

The bet & beachhead · **jobs to be done** · the riskiest assumptions

LAYER 2

## Test

### Run discovery

Get the truth, not politeness

Get out of the building · **the Mom Test** · what counts as evidence

LAYER 3

## Learn & re-aim

### Measure fit

Decide on evidence

The loop: move or hold · **measuring fit (the 40% test)** · the PMF decision

LAYER 1 · AIM

# Frame / the bet.

A defensible bet is a falsifiable hypothesis: if you can't say what would prove you wrong, you don't have a strategy — you have a guess you can't test. Frame it specifically enough that ten conversations could break it.

The bet & beachhead

Jobs to be Done

The riskiest assumptions



## TOOL 1.1 · THE BET &amp; BEACHHEAD

# The bet / & beachhead

*State your bet as one falsifiable sentence aimed at one narrow customer. “The space industry” is not a customer. The narrower and more specific the beachhead, the more testable — and winnable — the bet.*

## — WHAT IT IS

Your bet is a hypothesis about a **specific problem**, a **specific customer**, and why now. The customer must be a **beachhead** — one job title, one sector, one geography — not a market category. Geoffrey Moore’s logic: win the narrowest segment that feels the pain most acutely, earn its references, then expand (the bowling-pin strategy). Specificity is not a limitation; it’s what makes the bet testable in ten conversations and winnable with limited resources. “EO analysts at top-10 EU P&C carriers, DACH first” — not “insurers.”

## — THE QUESTIONS IT ANSWERS

- ◇ Can I state the bet as **one falsifiable sentence**?
- ◇ Is the beachhead **one title, one sector, one geography**?
- ◇ Who feels this pain **most acutely** and can act on it?
- ◇ If we’re right, what **unlocks** — what becomes possible?

## — WHAT "GOOD" LOOKS LIKE

- ✓ A single, specific, falsifiable problem statement
- ✓ A beachhead narrow enough to saturate
- ✓ The pain is acute, not a nice-to-have
- ✓ Specific enough that 10 calls could break it

## — WHAT NOT TO DO — AND WHY

**Aiming at a market instead of a beachhead.** “We’re for the insurance industry” feels ambitious but is untestable and unwinnable — you can’t saturate it, can’t get clean references, and can’t tell whether a “no” means the idea is wrong or just the wrong sub-segment. Vague problem statements (“space data is hard to access”) hide the real, specific pain. Narrow until it feels almost uncomfortably small — that’s where fit is found first.

Patrick Connolly

## — BEACHHEAD, NOT MARKET

TOO BROAD “The insurance market” — untestable

BEACHHEAD **One title + sector + geography**

Win the narrowest painful segment, earn references, then expand.

## ■ AEGIS E0

Aegis’s Day-1 bet was deliberately narrow and falsifiable: **“EO analysts at top-10 EU P&C carriers spend 4–6 hours per claim manually correlating satellite imagery with claim addresses.”** One role, one sector, DACH first. That specificity is exactly what let 22 analyst conversations *break* the first version of the bet within days — a vague bet would have survived on ambiguity.

## TOOL 1.2 · JOBS TO BE DONE

# Jobs / to be Done

*People don't buy products; they “hire” them to make progress in a specific situation. Frame the bet around the job the customer is trying to get done, not the features you want to build.*

## — WHAT IT IS

Jobs to be Done (Christensen) reframes demand: customers **hire** a product to do a **job** — to make progress in a particular circumstance. The job is stable; your solution is replaceable. Understanding it means understanding the **functional, emotional, and social** dimensions of the progress they want, and the **forces** at play: the push of the current pain and pull of the new way, against the anxiety of switching and the habit of today. Get the job right and the feature set follows; get it wrong and you build elegant answers to questions no one asked.

## — THE QUESTIONS IT ANSWERS

- ◇ What **job** is the customer hiring this to do?
- ◇ What do they **do today** to get that job done?
- ◇ What's the **push** (current pain) and **pull** (the new way)?
- ◇ What **anxiety or habit** keeps them from switching?

## — WHAT "GOOD" LOOKS LIKE

- ✓ The job is defined by the customer's progress, not your features
- ✓ You know the current alternative they 'hire' now
- ✓ Functional, emotional, and social dimensions surfaced
- ✓ Switching forces (push/pull vs. anxiety/habit) understood

## — WHAT NOT TO DO — AND WHY

**Falling in love with the solution, not the job.** Founders fixate on the clever thing they want to build and reverse-engineer a job to justify it. The tell: you can describe your features in detail but not the customer's actual workflow today. If you don't know what they currently 'hire' to do this job — a spreadsheet, a consultant, nothing at all — you don't yet understand the job, and your product is a guess.

Patrick Connolly

## — THE FORCES OF PROGRESS

|         |                           |
|---------|---------------------------|
| PUSH    | Pain with today's way     |
| PULL    | Appeal of the new way     |
| ANXIETY | Fear of switching         |
| HABIT   | Inertia of the status quo |

## ■ AEGIS E0

Aegis's first bet assumed the job was “get higher-resolution imagery.” Discovery revealed the real job: “**get the imagery I already pay for into my claims decision without 4–6 hours of manual work.**” Same customer, completely different job — and what they ‘hired’ today was a Maxar subscription sitting unused plus manual workarounds. Reframing around the actual job is what turned a dead bet into a live one.

## TOOL 1.3 · THE RISKIEST ASSUMPTIONS

# The riskiest / assumptions

*Every bet rests on a few assumptions that, if wrong, sink it. Name them explicitly, decide what evidence would falsify each, and aim your discovery straight at them.*

## — WHAT IT IS

A bet is only as strong as its hidden assumptions. List the **three riskiest** — the ones that, if false, mean the venture doesn't work — and for each write **exactly what would falsify it**. This is the discipline that separates testing from theatre: you decide the kill criterion *before* you go looking, so you can't rationalise away the evidence later. Then aim outreach at the assumptions, not at friendly contacts who'll confirm what you hope. Each assumption ends the sprint **confirmed, killed, or open**.

## — THE QUESTIONS IT ANSWERS

- ◇ What are the **3 assumptions** that, if wrong, sink the bet?
- ◇ For each, what **specific evidence would falsify it**?
- ◇ Am I aiming discovery at the **risky** assumptions, or the safe ones?
- ◇ Is each one ending up **confirmed, killed, or open**?

## — WHAT "GOOD" LOOKS LIKE

- ✓ The few make-or-break assumptions named explicitly
- ✓ A falsification test written for each, in advance
- ✓ Discovery aimed at the risk, not at reassurance
- ✓ Every assumption resolved to a clear status

## — WHAT NOT TO DO — AND WHY

**Testing the assumptions you're confident in, avoiding the scary ones.** Founders instinctively probe what they already believe and steer around the assumption most likely to be fatal — because confirming it feels good and breaking it hurts. But the riskiest assumption is exactly the one worth testing first and hardest. Setting the falsification criterion *before* the conversations is what stops you from quietly moving the goalposts

Patrick Connolly  
When the data is uncomfortable.

## — ASSUMPTION → STATUS

|           |                              |
|-----------|------------------------------|
| NAME IT   | The make-or-break assumption |
| FALSIFIER | What would prove it wrong    |
| TEST      | Aim discovery at it          |
| STATUS    | Confirmed / killed / open    |

## ■ AEGIS E0

Aegis named three make-or-break assumptions up front, each with a falsifier. "Analysts have €50k vendor authority" — **killed** (11 of 13 said authority sits with Group IT). "Analysts spend 4–6 hrs/claim correlating imagery" — **killed** (16 of 22 spent <1 hr or skipped it). "No competitor solves integration for EU mid-market" — **confirmed**. Writing the falsifiers first is why these resolved cleanly instead of being argued away.

LAYER 2 • TEST

# Run / discovery.

Everyone tells you to talk to customers; almost no one tells you how to do it without being misled. Discovery done wrong collects compliments. Done right, it surfaces the truth — and the truth is in what people have actually done, not what they say they'd do.

Get out of the building

The Mom Test

What counts as evidence

2

## TOOL 2.1 · GET OUT OF THE BUILDING

# Get out / of the building

*No amount of analysis inside the office substitutes for talking to real buyers. Assumptions are to be tested, not trusted — and the only way to test them is to get out and have real conversations, at volume.*

## — WHAT IT IS

Steve Blank's customer-development principle: the facts you need are **outside the building**, with real customers, not in your own conviction or a LinkedIn poll. Treat every conversation as **data collection**, not selling. Go for **volume across segments** — buyers, users, and adjacent players, not 50 of the same role — because patterns only appear in a real dataset. Log every conversation **same-day, in their words**: the tracker is the evidence base your decisions (and your investors' trust) rest on. If it's not written down, it didn't happen.

## — THE QUESTIONS IT ANSWERS

- ◇ *Am I testing assumptions with **real buyers**, not my team?*
- ◇ *Am I talking to **enough** people to see a pattern (aim ~50)?*
- ◇ *Across **multiple segments** — users, buyers, adjacent?*
- ◇ *Is every conversation **logged same-day, in their words**?*

## — WHAT "GOOD" LOOKS LIKE

- ✓ Conversations with real buyers, at volume
- ✓ Multiple segments, not 50 of one role
- ✓ Warm intros used (they convert 5–10x cold)
- ✓ Same-day logging, quotes captured, surprises noted

## — WHAT NOT TO DO — AND WHY

**Substituting conviction for conversations.** The most common failure is building from inside the building — trusting your own certainty, your team's agreement, or a few friendly replies, instead of doing the unglamorous work of dozens of real conversations. Talking only to people who already like you, or logging "sent a LinkedIn message" as a

Patrick Connolly

## — THE DISCOVERY DATASET

|        |                           |
|--------|---------------------------|
| VOLUME | ~50 real conversations    |
| SPREAD | Users + buyers + adjacent |
| LOG    | Same-day, in their words  |
| RESULT | Patterns you can trust    |

## ■ AEGIS E0

Aegis ran **52 real conversations across 4 segments** in ten days — 22 analysts/underwriters, 13 buyers (Group IT / procurement), 10 brokers, 7 adjacent (Maxar, consultants). Every call logged same-day in the customer's words. That dataset is what made the patterns undeniable: **the bet shifted three times on evidence**, not on any single persuasive call.



## TOOL 2.2 · THE MOM TEST

# The / Mom Test

*Even your mum will lie to protect your feelings. The fix isn't better people to ask — it's better questions: talk about their life and their past, never your idea and their future.*

## — WHAT IT IS

Rob Fitzpatrick's rules for questions that survive bias: (1) talk about **their life**, not your idea; (2) ask about **specifics in the past**, not generics or hypotheticals about the future; (3) **talk less, listen more**. Future-tense questions ("Would you use a tool that...?") generate polite agreement — a known false positive. Past-behaviour questions ("Walk me through the last time you faced this") generate facts: what they did, what it cost, what they've already tried to fix it. Crucially, **don't pitch** — the moment you pitch, you stop learning.

## — THE QUESTIONS IT ANSWERS

- ◇ *Am I asking about **their life**, or pitching **my idea**?*
- ◇ *Am I asking about the **past** ("last time you..."), not the future?*
- ◇ *Have I found where they **already spend time and money** on this?*
- ◇ *Am I **listening** more than I'm talking?*

## — WHAT "GOOD" LOOKS LIKE

- ✓ Questions about real past behaviour, not opinions
- ✓ No pitching — the idea stays out of the conversation
- ✓ Evidence of existing spend / effort on the problem
- ✓ Compliments treated as warning signs, not validation

## — WHAT NOT TO DO — AND WHY

**Collecting compliments instead of facts.** "Would you use this?" → "Sure!" feels like validation and means nothing — you fed them the answer. "Do you like my idea?" invites politeness. These false positives cost founders months of building the wrong thing. The discipline is uncomfortable: ask the questions that could expose your idea's flaws, chase history not hypotheticals, and never let them be nice to you instead of honest.

Patrick Connolly

## — ASK THIS, NOT THAT

- |     |                                                     |
|-----|-----------------------------------------------------|
| NOT | "Would you use a tool that...?"                     |
| ASK | "Walk me through the last time you faced this."     |
| NOT | "How much would you pay?"                           |
| ASK | "How are you solving it today — what's it costing?" |

## ■ AEGIS E0

Aegis's playbook was explicitly "**ask, don't pitch**." Instead of "would you buy an integration layer?" they asked analysts to **walk through their last claim involving imagery** — which surfaced, unprompted, that they pay for Maxar and can't use it. The strongest signal of all came as past-tense action, not opinion: Munich Re **signed a pilot MOU**. "Send me the agreement" is a yes; "sounds useful, good luck" is a polite no.

## TOOL 2.3 · WHAT COUNTS AS EVIDENCE

# What counts / as evidence

*Not all signals are equal. Before you let a conversation move your bet, test it against four bars — or you'll pivot on noise and ignore the patterns that matter.*

## — WHAT IT IS

A signal earns the right to change your strategy only when it clears four tests.

**Pattern, not one-off:** many people, not one enthusiast. **Unprompted, not led:** they named the pain before you mentioned it. **Buyer, not bystander:** it came from someone who'd actually hold the budget, not a friendly observer. **Action, not opinion:** “send me the pilot agreement” outweighs “sounds useful.” Weigh evidence on these axes and the dataset tells you what's real — distinguishing a genuine pull signal from flattering noise.

## — THE QUESTIONS IT ANSWERS

- ◇ Is this a **pattern**, or one persuasive person?
- ◇ Did they raise it **unprompted**, or did I lead them?
- ◇ Is it from a **real buyer**, or a bystander?
- ◇ Is it an **action** (“send the contract”), or just an opinion?

## — WHAT "GOOD" LOOKS LIKE

- ✓ Conclusions drawn from patterns, not anecdotes
- ✓ Unprompted pain weighted above prompted agreement
- ✓ Buyer signals separated from bystander interest
- ✓ Actions (pilots, intros, payment) treated as the real yes

## — WHAT NOT TO DO — AND WHY

**Treating a flattering anecdote as proof.** One excited person at a party, a led “yes,” or a friendly analyst with no budget can feel like validation and send you down a months-long wrong path. The mirror error is ignoring an unprompted pattern from real buyers because it contradicts the bet. Score every signal on the four bars before acting; a strong pull signal is unprompted, patterned, from a buyer, and attached to an action.

Patrick Connolly

## — THE FOUR-BAR TEST

- 1 Pattern, not one-off
- 2 Unprompted, not led
- 3 Buyer, not bystander
- 4 Action, not opinion

## ■ AEGIS E0

Aegis's decisions were tied to signals that cleared all four bars: “we pay for imagery, can't use it” was a **pattern** (14 of 22), **unprompted**, from **analysts and confirmed by buyers**, and ultimately backed by **action** — Munich Re's signed MOU and Allianz referring two peer carriers. A single enthusiastic founder-friend would have been logged and discounted as noise.

LAYER 3 · LEARN & RE-AIM

# Measure / fit.

Strategy and discovery are one engine: the bet aims your outreach, evidence comes back, and you decide — move the bet or hold it. Then measure fit with a real leading indicator, and make the call on data, not nerves.

The loop: move or hold

Measuring fit (the 40% test)

The PMF decision

3

## TOOL 3.1 · THE LOOP: MOVE OR HOLD

# The loop / move or hold

*Every conversation should confirm the bet, sharpen it, or break it. The skill is knowing when evidence should move your strategy — and when holding steady is the disciplined choice.*

## — WHAT IT IS

Strategy and outreach run as one loop: **aim** (the bet) → **test** (outreach as data) → **learn** (what came back) → **re-aim** (move or hold), and round again. Build–Measure–Learn (Ries) is the same engine. The judgment is when to **pivot vs. persevere**: change the bet only on *real* evidence — a pattern, unprompted, from a buyer, attached to action. Otherwise hold and keep testing. A pivot is a change in strategy with a clear reason serving a coherent vision — not a reaction to the last conversation or to your own nerves.

## — THE QUESTIONS IT ANSWERS

- ◇ Did this round of evidence **confirm, sharpen, or break** the bet?
- ◇ Is the signal strong enough to **move** the bet (all four bars)?
- ◇ Am I pivoting on **learning**, or on anxiety?
- ◇ If I hold, am I giving the bet a **fair test** first?

## — WHAT "GOOD" LOOKS LIKE

- ✓ Each conversation classified: confirm / sharpen / break
- ✓ Bet moves only on evidence that cleared the four bars
- ✓ Pivots have a clear, stated reason
- ✓ Holds are deliberate, not avoidance of bad news

## — WHAT NOT TO DO — AND WHY

**Two failure modes that both feel like progress.** The *Over-Pivoter* changes the bet after every surprising call — never runs a test long enough to learn, so nothing is ever proven. The *Confirmer* hears only what supports the bet, logs every call as “confirmed,” and ends day ten believing exactly what they believed on day one — 50 conversations, zero learning. Don’t pivot because you’re nervous; pivot because you learned something real.

Patrick Connolly

## — THE LOOP



Avoid both: the Over-Pivoter (thrash) and the Confirmer (theatre).

## ■ AEGIS E0

Aegis ran the loop honestly. The Day-1 bet — “insurers will pay for higher-res imagery” — broke on a clear pattern (“we pay for Maxar, it sits unused,” x14, unprompted), so they **pivoted to integration**. When later calls were merely flattering, they **held**. Three real moves over ten days, each tied to evidence that cleared the bars — not thrashing, not theatre.

## TOOL 3.2 · MEASURING FIT

# Measuring fit / the 40% test

*PMF feels obvious in hindsight but is hard to see in the moment. Sean Ellis found a leading indicator you can measure before the growth shows up — and an engine to improve it deliberately.*

## — WHAT IT IS

Once people are using the product, measure fit with a **leading** indicator, not a lagging one. Sean Ellis's test: ask users **“how would you feel if you could no longer use this?”** and track the percent who say **“very disappointed.”** Across ~100 startups, **40%** separated those that struggled to grow from those that grew easily. Rahul Vohra's **Superhuman engine** makes it actionable: (1) **segment** to the users who'd be very disappointed, (2) **analyse** why they love it and what holds the “somewhat” group back, (3) build a **50/50 roadmap** (deepen the love / remove blockers), (4) **repeat**. Superhuman went 22% → 58%.

## — THE QUESTIONS IT ANSWERS

- ◇ What % of users would be **“very disappointed”** without us?
- ◇ Who are the users that **love it** — and what do they share?
- ◇ What holds the **“somewhat disappointed”** group back?
- ◇ Is our roadmap **50/50** — deepen love, remove blockers?

## — WHAT "GOOD" LOOKS LIKE

- ✓ A measured PMF score, not a gut feeling
- ✓ The high-expectation customer profile identified
- ✓ Roadmap split between deepening love and fixing blockers
- ✓ Score re-measured on a cadence — it's the key metric

## — WHAT NOT TO DO — AND WHY

**Waiting for lagging indicators — or measuring everyone.** Revenue and press are lagging: by the time they confirm PMF, you already had it. The 40% test is a leading signal you can act on early. But measure the *right* users (those who've genuinely experienced the core product, e.g. used it 2+ times recently), and don't average across everyone — a

Patrick Connolly

## — THE SUPERHUMAN ENGINE

Survey (“very disappointed?”) → Segment → 50/50 roadmap →

Repeat

40% “very disappointed” = initial PMF.  
Superhuman: 22% → 58%.

## ■ AEGIS E0

As Aegis moved from discovery to a working concierge MVP at Munich Re, the 40% test became its north star. Early on, only the underwriting-data team would have been “very disappointed” — so rather than chase a broad average, Aegis **segmented to that high-expectation customer** and built the roadmap around what they loved (time saved) and what blocked the rest (PII assurance). That's the engine: measure, segment, improve, repeat.

Founder Toolkit

## TOOL 3.3 · THE PMF DECISION

# The PMF / decision

*Fit isn't one survey number — it's a convergence of signals. Know what genuine fit looks like, decide what you'll do at each outcome, and don't declare victory (or defeat) prematurely.*

## — WHAT IT IS

Triangulate fit across signals, not one metric: the **40% test** clearing the bar; **retention curves that flatten** (a cohort that sticks rather than bleeding to zero — the single strongest behavioural signal); **organic pull** (referrals, unprompted word of mouth, inbound); and real **buyer action** (pilots, paid contracts). Decide your response in advance: strong convergence → start building the repeatable go-to-market motion (the next toolkit); mixed → segment to where fit is real and deepen it; weak → return to the bet. Fit is the milestone that earns the right to scale.

## — THE QUESTIONS IT ANSWERS

- ◇ Do **multiple signals converge** — survey, retention, pull, action?
- ◇ Does a cohort's **retention curve flatten**, or bleed to zero?
- ◇ Is growth becoming **organic** (referrals, inbound)?
- ◇ What will we **do** at strong / mixed / weak fit — decided in advance?

## — WHAT "GOOD" LOOKS LIKE

- ✓ Fit judged on converging signals, not one number
- ✓ Retention curve flattening confirmed in a real cohort
- ✓ Organic pull emerging, not just paid push
- ✓ A pre-committed response to each fit outcome

## — WHAT NOT TO DO — AND WHY

**Declaring PMF on a single flattering signal — then scaling.** A good survey score with retention bleeding to zero is not fit; pouring money into growth on top of a leaky product just burns cash faster (the premature-scaling trap). The opposite error is endless tinkering when fit is already there. Name the converging signals you'd need, decide the response to each outcome before you measure, and only scale once fit is real.

Patrick Connolly

## — SIGNALS THAT CONVERGE

|           |                                   |
|-----------|-----------------------------------|
| SURVEY    | 40%+ “very disappointed”          |
| RETENTION | Cohort curve flattens             |
| PULL      | Referrals, inbound, word of mouth |
| ACTION    | Pilots, paid contracts            |

## ■ AEGIS E0

For Aegis, the convergence was clear enough to act on: the high-expectation segment would be very disappointed without the integration, the pilot cohort **kept using it**, and pull was organic — **Allianz referred two peer carriers unprompted** and AXA requested follow-up. That convergence — not any single number — is what earned Aegis the right to move from finding fit to **building a repeatable go-to-market motion**.

PUTTING IT TOGETHER

# Find the fit / before you build it.

Product-market fit is earned through evidence, not conviction. Frame a falsifiable bet, run discovery that gets the truth, then measure fit honestly and decide. The whole game is one loop — aim, test, learn, re-aim — run on purpose until the signals converge. Then, and only then, you've earned the right to scale.

Layer 1

## Aim

A falsifiable bet, a narrow beachhead, the real job, and the riskiest assumptions named with their kill criteria.

Layer 2

## Test

Get out of the building, ask about the past not the future, and weigh evidence on the four bars.

Layer 3

## Learn

Run the loop, measure fit with the 40% test, and decide on converging signals — not nerves.

Next

## Then scale

Real, converged fit earns the right to build a repeatable go-to-market motion. That's the next toolkit.



## THE EVIDENCE BASE

# Sources / & further reading.

This toolkit builds on the ISU PMF Sprint and grounds it in the established customer-development and product-market-fit literature. Much of this field is rigorous practitioner work rather than peer-reviewed research; where genuine empirical or academic grounding exists (Jobs to be Done, customer development, the 40% benchmark), it is cited. Aegis EO is a fictional teaching venture; its sprint mirrors the worked example in the ISU tracker.

---

## Customer Development

Steve Blank, *The Four Steps to the Epiphany*. “Get out of the building”; assumptions are to be tested, not trusted; discovery → validation.

---

## The Mom Test (2013)

Rob Fitzpatrick. Talk about their life not your idea; ask about the past not the future; don’t pitch. The discovery-interview standard.

---

## The PMF survey & 40% benchmark

Sean Ellis. “How would you feel if you could no longer use this?” — 40% “very disappointed” as a leading indicator, benchmarked across ~100 startups.

---

## Beachhead / bowling-pin strategy

Geoffrey Moore, *Crossing the Chasm*. Win the narrowest painful segment, earn references, then expand.

---

## Retention as the fit signal

Industry consensus (Reforge / a16z): a cohort retention curve that flattens is the strongest behavioural evidence of fit.

---

## The Lean Startup (2011)

Eric Ries. Build–Measure–Learn, validated learning, pivot-or-persevere — the engine behind the loop.

---

## Jobs to be Done

Clayton Christensen (HBS). Customers ‘hire’ products to make progress; the functional/emotional/social job and the forces of progress.

---

## The Superhuman PMF engine

Rahul Vohra (First Round Review). Segment → analyse → 50/50 roadmap → repeat; Superhuman went 22% → 58%.

---

## “The only thing that matters”

Marc Andreessen. PMF as the dividing line; “you can always feel when it’s happening” — a vivid but lagging definition.

---

## The ISU PMF Sprint

Patrick Connolly & Prof. J. Gaillard (ISU MSS26): the bet, the loop, the two failure modes, the four-bar evidence test, the tracker.